



# Introduction to DSP

Елена Кантонистова

По мотивам материалов Александра Марковича





# What is sound?

- **Sound wave** is the pattern of **oscillations** caused by the movement of energy traveling through the air
- **Microphone** picks up these air **oscillations** and converts them into electrical vibrations
- These **oscillations** are converted into an **analog** signal and then a **digital** signal




# How is sound stored in the computer?

- The **analog** signal is discretized, quantized and encoded
- An analog signal is **discretized** in that the signal is represented as a sequence of values taken at discrete points in time  $t$  with step  $d$
- **Quantisation** of a signal consists in splitting the range of signal values into  $N$  levels in increments of  $d$  and selecting for each reference the level that corresponds to it
- Signal **encoding** is just a way of presenting the signal in a more compact form



# Analog-to-Digital Conversion


- Converting analog signals to a sequence of numbers having finite precision
- Corresponding devices are called A/D converters (ADCs)



# What other characteristics are there?

- **Sample rate (SR)** - number of audio samples per one second (e.g. 8 kHz, 22.05 kHz, 44.1 kHz)
- **Sample size** - number of bits per one sample (e.g. 8, 16, 25, 32 bits)
- **Number of channels** -- how many signals we record in parallel (e.g. mono(1), stereo(2))

## 8000 Hz

The international [G.711](#)  standard for audio used in telephony uses a sample rate of 8000 Hz (8 kHz). This is enough for human speech to be comprehensible.

## 44100 Hz

The 44.1 kHz sample rate is used for compact disc (CD) audio. CDs provide uncompressed 16-bit stereo sound at 44.1 kHz. Computer audio also frequently uses this frequency by default.

## 48000 Hz

The audio on DVD is recorded at 48 kHz. This is also often used for computer audio.

## 96000 Hz


High-resolution audio.

## 192000 Hz

Ultra-high resolution audio. Not commonly used yet, but this will change over time.


# What other characteristics are there?

- Assume  $\mathbf{f}(\mathbf{n})$  is our signal where  $\mathbf{n}$  is time
- Power of signal is  $f^2(n)$
- Energy of signal is  $\sum f^2(n)$
- In practice estimated by some **window**
- Energy in **decibels**:  $10 \log_{10} E$
- $\text{SNR}_{dB} = 10 \log_{10} \frac{E_{\text{signal}}}{E_{\text{noise}}}$



# What about audio formats?

- Non-compressed formats: **WAV, AIFF, etc.**
- Lossless compression(2:1) : **FLAC, ALAC, etc.**
- Lossy compression(10:1) : **MP3, Opus, etc**
- **Bit rate** measure a degree of compression. Number of bit that are conveyed or processed per **unit of time**.




# Fourier Transform

Complex sound wave




Find  
Fourier Series  
→

Combination of simple sound waves





# Fourier Transform




# Fourier Transform

- The **Fourier transform(FT)** is a mathematical formula that allows us to decompose a signal into its individual **frequencies** and the frequency's **amplitude**
- FT transfer a signal from the **time domain** to the **frequency domain**

- $$F(y) = \int_{-\infty}^{\infty} f(x) e^{-2\pi i x y} dx$$

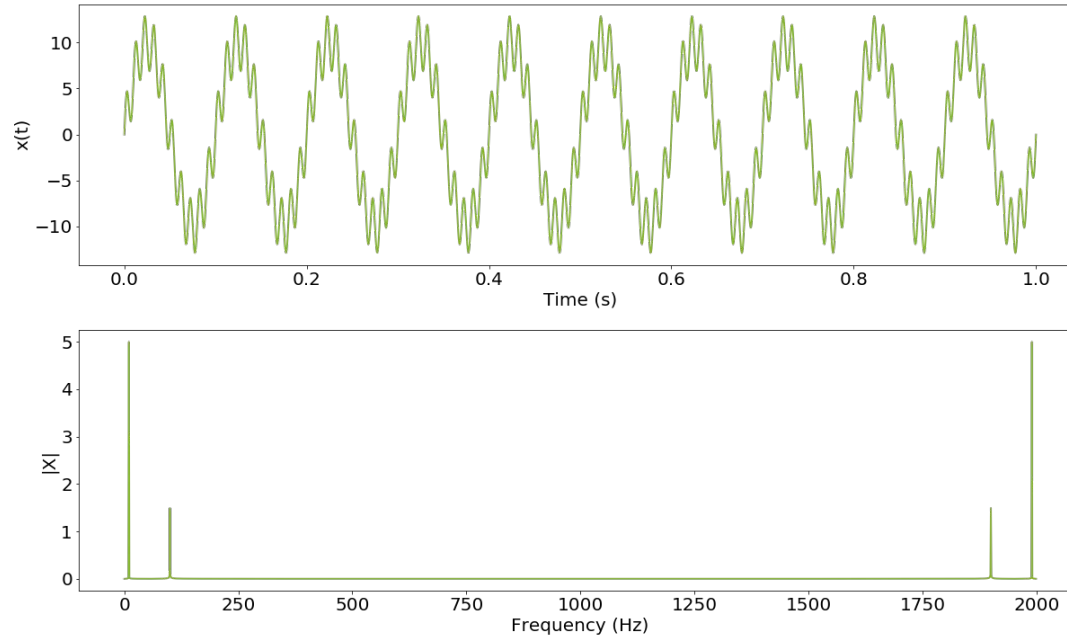
time  $\rightarrow$  frequency



# Example of DFT

$$F = 2kH z$$


$$f(t) = 10 \sin(2\pi 10t) + 3 \sin(2\pi 100t)$$




# Short-time Fourier transform

FFT + Windowing


Sliced signal



Window




Windowed signal




# Short-time Fourier transform

FFT + Windowing






# Spectrograms




# Mel Scale

$$m = 2595 \log_{10} \left( 1 + \frac{f}{700} \right) = 1127 \ln \left( 1 + \frac{f}{700} \right)$$

$$f = 700 \left( 10^{\frac{m}{2595}} - 1 \right) = 700 \left( e^{\frac{m}{1127}} - 1 \right)$$



# CNN-approach





# Automatic Speech Recognition

Task:

- Transform speech (audio) to text

Also known as:

- STT - Speech To Text



# WER

## Word Error Rate

### How to compute?

Edit path from reference to prediction

- **S** – substitution count
- **D** – deletions count
- **I** – insertions count
- **C** – correct count
- **N** – **S + D + C** – Total word count in reference

True: quick brown fox jumped over a lazy dog  
Pred: quick brow an fox jumped over lazy dog

$$WER = \frac{S + D + I}{N} = \frac{S + D + I}{S + D + C}$$

# CER

Character Error Rate

The same, but on character level

Questions:

- 1) what is more important?
- 2) what is more difficult to minimize?

# LibriSpeech

Domain:

- audiobooks

Parts:

- clean - low-WER speakers
- other - high-WER speakers

Features:

- 10-20s audio
- long sentences
- complex language

subset	hours	per-spkr minutes	female spkrs	male spkrs	total spkrs
dev-clean	5.4	8	20	20	40
test-clean	5.4	8	20	20	40
dev-other	5.3	10	16	17	33
test-other	5.1	10	17	16	33
train-clean-100	100.6	25	125	126	251
train-clean-360	363.6	25	439	482	921
train-other-500	496.7	30	564	602	1166

[1]

Human WER:

- test-clean: **5.83**
- test-other: **12.69**

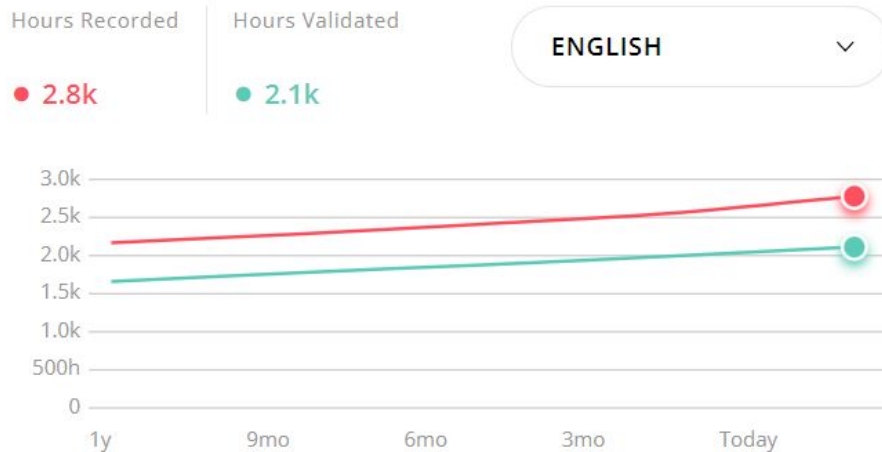
# Mozilla Common Voice

## Domain:

- Short random phrases
- Multiple languages

## Features:

- Crowdsourced
- Simple language
- Short phrases
- Frequently updated and validated


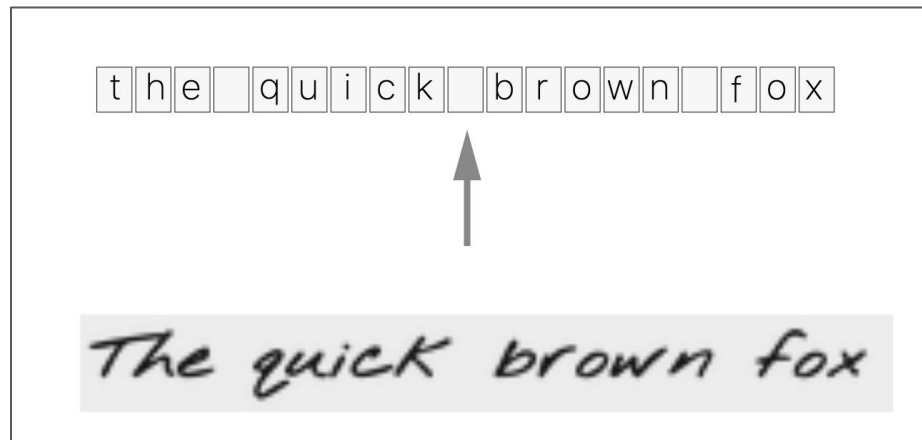


# Fundamental problem

## Alignment problem

- Variable length input
- Variable length output
- No alignment

## How to train?



[3]

# Connectionist Temporal Classification

2006

Idea:

- Split input into frames
- Classify each frame into *num\_letters* classes
- Merge consecutive letters

$x_1$   $x_2$   $x_3$   $x_4$   $x_5$   $x_6$

input ( $X$ )

c c a a a t

alignment

c a t

output ( $Y$ )

[3]

Issues?

# Connectionist Temporal Classification

Idea:

- Split input into frames
- Classify each frame into *num\_letters* classes
- Merge consecutive letters

$x_1$   $x_2$   $x_3$   $x_4$   $x_5$   $x_6$

input ( $X$ )

c c a a a t

alignment

c a t

output ( $Y$ )

[3]

Issues?

- Multiple consecutive letters in target word (e.g.: hello)
- Silence between words and letters

# CTC: Empty token

h h e  $\epsilon$   $\epsilon$  | | |  $\epsilon$  | | o

h e  $\epsilon$  |  $\epsilon$  | o

h e | | o

h e | | o

[3]

First, merge repeat characters.

Then, remove any  $\epsilon$  tokens.

The remaining characters are the output.



# CTC: Empty token

## Valid Alignments

€ c c € a t

c c a a t t

c a € € € t

[3]

## Invalid Alignments

c € c € a t


corresponds to  
 $Y = [c, c, a, t]$

c c a a t   


has length 5

c € € € | t t


missing the 'a'




We start with an input sequence,  
like a spectrogram of audio.



The input is fed into an RNN,  
for example.




We start with an input sequence,  
like a spectrogram of audio.




The input is fed into an RNN,  
for example.

h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
€	€	€	€	€	€	€	€	€	€

The network gives  $p_t(a | X)$ ,  
a distribution over the outputs  
 $\{h, e, l, o, \epsilon\}$  for each input step.



We start with an input sequence,  
like a spectrogram of audio.




The input is fed into an RNN,  
for example.

h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
€	€	€	€	€	€	€	€	€	€


The network gives  $p_t(a | X)$ ,  
a distribution over the outputs  
 $\{h, e, l, o, \epsilon\}$  for each input step.

h	e	€	l	l	€	l	l	o	o
h	h	e	l	l	€	€	l	€	o
€	e	€	l	l	€	€	l	o	o


With the per time-step output  
distribution, we compute the  
probability of different sequences




We start with an input sequence, like a spectrogram of audio.




The input is fed into an RNN, for example.



The network gives  $p_t(a | X)$ , a distribution over the outputs  $\{h, e, l, o, \epsilon\}$  for each input step.





With the per time-step output distribution, we compute the probability of different sequences



By marginalizing over alignments, we get a distribution over outputs.

[3]



h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
€	€	€	€	€	€	€	€	€	€

h	e	€	l	l	€	l	l	o	o
h	h	e	l	l	€	€	l	€	o
€	e	€	l	l	€	€	l	o	o

h	e	l	l	o
e	l	l	o	
h	e	l	o	

[3]

Loss Function:

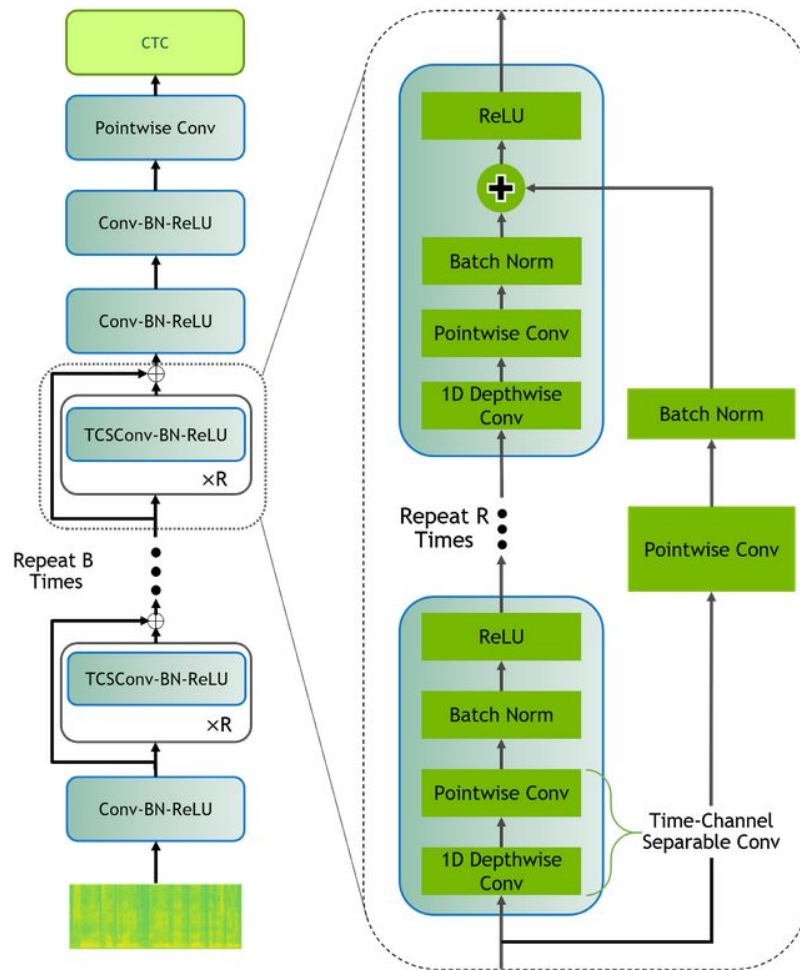
$$p(Y | X) = \sum_{A \in \mathcal{A}_{X,Y}} \prod_{t=1}^T p_t(a_t | X)$$

The CTC conditional **probability**

**marginalizes** over the set of valid alignments

computing the **probability** for a single alignment step-by-step.

# QuartzNet



# Language Models (LM): motivation

## Problem:

- Spelling of a word heavily depends on its context
- We always want **more** data
- Labeled audio data is difficult to obtain

hypo 1: let's go **two** a movie (score: 0.21)  
hypo 2: let's go **to** a movie (score: 0.19)  
hypo 3: let's go **too** a movie (score: 0.13)

## Idea:

- ASR predicts texts
- Unlabeled text data is **very easy** to get
- Let's make use of it!



# LM: motivation

**Language model** - a model that estimates the probability of a text.

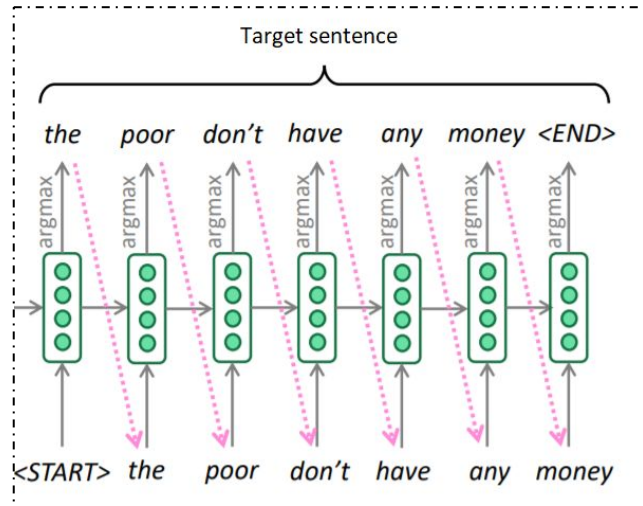
Usually requires unlabeled text corpus to train.

Examples:

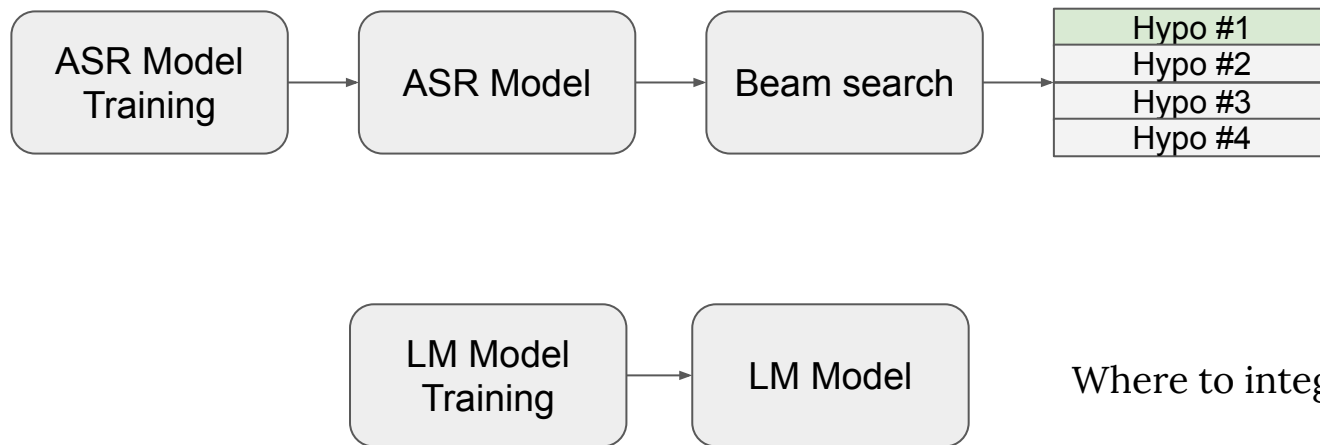
- N-gramms (simple)
- Neural networks (complex) e.g: Bert, GPT-3

$P(\text{let's go } \mathbf{two} \text{ a movie}) = 0.01$

$P(\text{let's go } \mathbf{to} \text{ a movie}) = 0.6$



# LM: ASR pipeline



# LM: final hypothesis rescoring

a.k.a.: *second-pass rescoring*

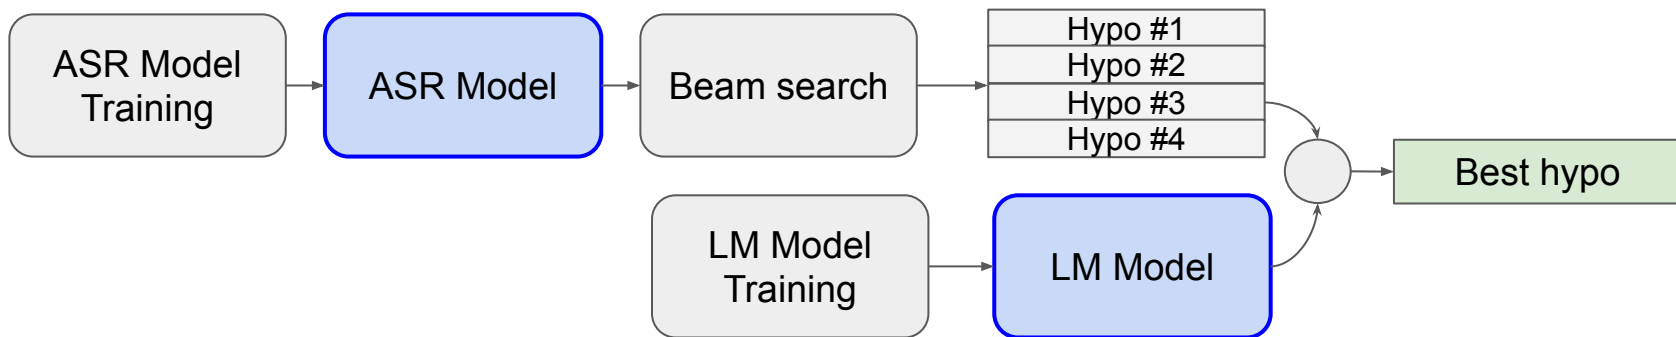
Prerequisites:

- Trained ASR
- Trained LM

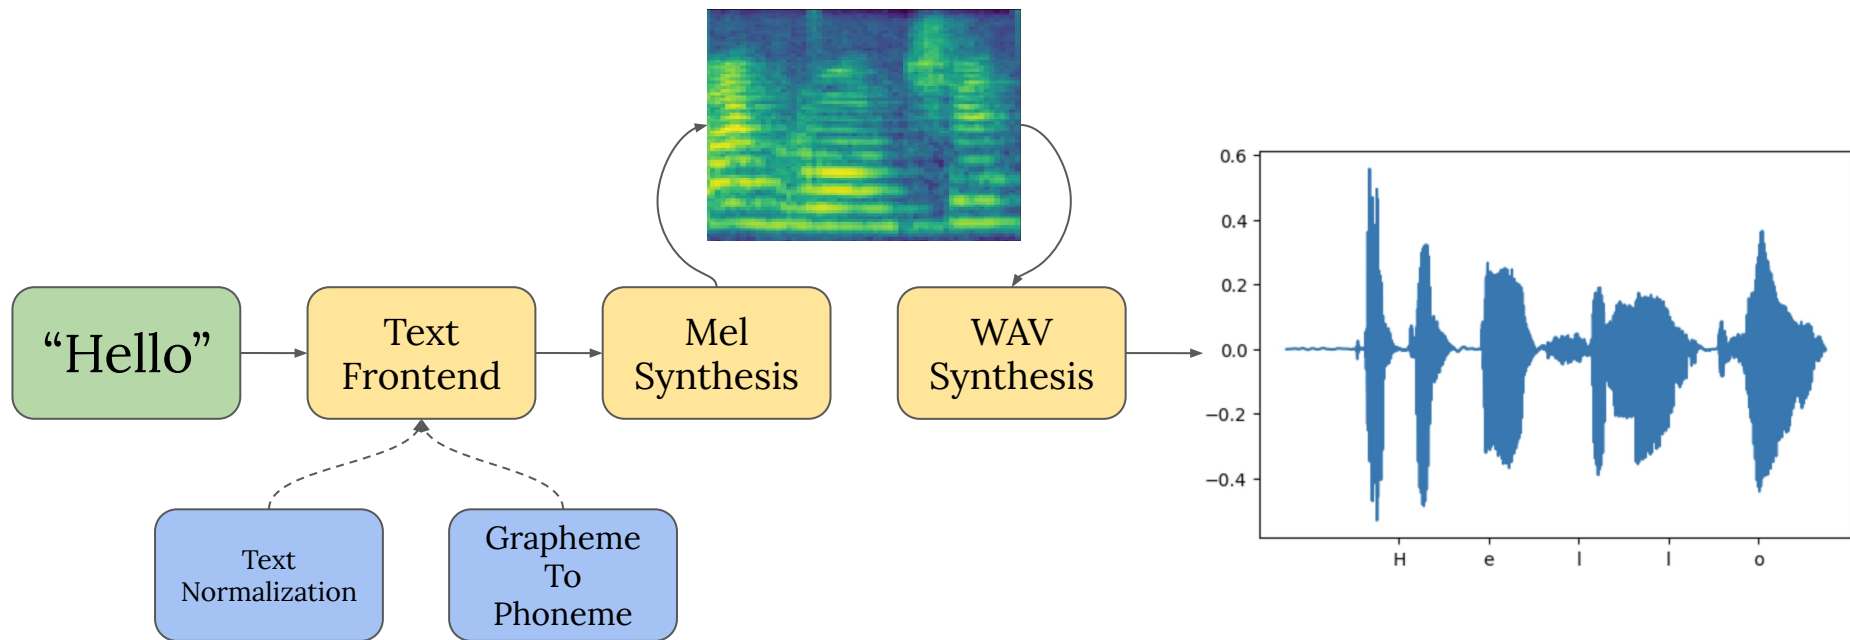
Concept: Rescore beam-search output with LM scores

$$P_{final} = P_{ASR}(Y|X) + \alpha \cdot P_{LM}(Y) + \beta \cdot len(Y)$$

**Oracle WER** – lowest WER among all beam search output hypothesis



# Text-to-Speech (TTS)



# Tacotron 2

